

Neural Network Basics

LATEST SUBMISSION GRADE

100%

1. What does a neuron compute?

1 / 1 point

- ☐ A neuron computes an activation function followed by a linear function ($z = Wx + b$)
- ☐ A neuron computes the mean of all features before applying the output to an activation function
- ☐ A neuron computes a function g that scales the input x linearly ($Wx + b$)
- ☒ A neuron computes a linear function ($z = Wx + b$) followed by an activation function

✓ Correct

Correct, we generally say that the output of a neuron is $a = g(Wx + b)$ where g is the activation function (sigmoid, tanh, ReLU, ...).

2. Which of these is the "Logistic Loss"?

1 / 1 point

- ☒ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = -(y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$
- ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = \max(0, y^{(i)} - \hat{y}^{(i)})$
- ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = |y^{(i)} - \hat{y}^{(i)}|^2$
- ☐ $\mathcal{L}^{(i)}(\hat{y}^{(i)}, y^{(i)}) = |y^{(i)} - \hat{y}^{(i)}|$

✓ Correct

Correct, this is the logistic loss you've seen in lecture!

3. Suppose `img` is a (32,32,3) array, representing a 32x32 image with 3 color channels red, green and blue. How do you reshape this into a column vector?

1 / 1 point

- ☐ `x = img.reshape((32*32,3))`
- ☐ `x = img.reshape((3,32*32))`
- ☐ `x = img.reshape((1,32*32,*3))`
- ☒ `x = img.reshape((32*32*3,1))`

✓ Correct

4. Consider the two following random arrays "a" and "b":

1 / 1 point

```
1 a = np.random.randn(2, 3) # a.shape = (2, 3)
2 b = np.random.randn(2, 1) # b.shape = (2, 1)
3 c = a + b
```

What will be the shape of "c"?

- ☐ c.shape = (2, 1)
- ☐ c.shape = (3, 2)
- ☒ c.shape = (2, 3)
- ☐ The computation cannot happen because the sizes don't match. It's going to be "Error"!

✓ Correct

Yes! This is broadcasting. b (column vector) is copied 3 times so that it can be summed to each column of a.

5. Consider the two following random arrays "a" and "b":

1 / 1 point

```
1 a = np.random.randn(4, 3) # a.shape = (4, 3)
2 b = np.random.randn(3, 2) # b.shape = (3, 2)
3 c = a*b
```

What will be the shape of "c"?

- ☐ c.shape = (3, 3)
- ☐ c.shape = (4, 3)
- ☒ The computation cannot happen because the sizes don't match. It's going to be "Error"!
- ☐ c.shape = (4,2)

✓ Correct

Indeed! In numpy the "*" operator indicates element-wise multiplication. It is different from "np.dot()". If you would try "c = np.dot(a,b)" you would get c.shape = (4, 2).

6. Suppose you have n_x input features per example. Recall that $X = [x^{(1)} x^{(2)} \dots x^{(m)}]$. What is the dimension of X ?

1 / 1 point

- ☐ $(m, 1)$
- ☐ (m, n_x)
- ☐ $(1, m)$
- ☒ (n_x, m)

✓ Correct

7. Recall that "`np.dot(a,b)`" performs a matrix multiplication on a and b , whereas "`a*b`" performs an element-wise multiplication.

1 / 1 point

Consider the two following random arrays " a " and " b ":

```
1 a = np.random.randn(12288, 150) # a.shape = (12288, 150)
2 b = np.random.randn(150, 45) # b.shape = (150, 45)
3 c = np.dot(a,b)
```

What is the shape of c ?

- ☐ $c.shape = (12288, 150)$
- ☒ $c.shape = (12288, 45)$
- ☐ The computation cannot happen because the sizes don't match. It's going to be "Error!"
- ☐ $c.shape = (150, 150)$

✓ Correct

Correct, remember that a `np.dot(a, b)` has shape (number of rows of a , number of columns of b). The sizes match because :

"number of columns of $a = 150 =$ number of rows of b "

8. Consider the following code snippet:

1 / 1 point

```
1 # a.shape = (3,4)
2 # b.shape = (4,1)
3
4 for i in range(3):
5     for j in range(4):
6         c[i][j] = a[i][j] + b[j]
```

How do you vectorize this?

- ☐ $c = a.T + b$
- ☐ $c = a + b$
- ☐ $c = a.T + b.T$
- ☒ $c = a + b.T$

✓ Correct

9. Consider the following code:

1 / 1 point

```
1 a = np.random.randn(3, 3)
2 b = np.random.randn(3, 1)
3 c = a*b
```


What will be c? (If you're not sure, feel free to run this in python to find out).

- ☒ This will invoke broadcasting, so b is copied three times to become (3,3), and * is an element-wise product so c.shape will be (3, 3)
- ☐ This will invoke broadcasting, so b is copied three times to become (3, 3), and * invokes a matrix multiplication operation of two 3x3 matrices so c.shape will be (3, 3)
- ☐ This will multiply a 3x3 matrix a with a 3x1 vector, thus resulting in a 3x1 vector. That is, c.shape = (3,1).
- ☐ It will lead to an error since you cannot use "*" to operate on these two matrices. You need to instead use np.dot(a,b)

✓ Correct

10. Consider the following computation graph.

1 / 1 point



What is the output J?

- ☐ $J = (c - 1) * (b + a)$
- ☒ $J = (a - 1) * (b + c)$
- ☐ $J = a*b + b*c + a*c$
- ☐ $J = (b - 1) * (c + a)$

✓ Correct

Yes. $J = u + v - w = a*b + a*c - (b + c) = a * (b + c) - (b + c) = (a - 1) * (b + c)$.